

Unified Context Engine (UCE)

Explanation-First Technical Research Report

Vanilla LLM vs MCP-Powered UCE for Impact Analysis, Governance Tracing, and RBAC Control

Generated 2026-05-05 17:43:33 from repository artifacts in [research/icml2026/results](#)

Executive Summary

This report explains the Unified Context Engine (UCE) as a governance-aware context layer for agentic software engineering. The core idea is simple: a vanilla LLM can read text and make plausible guesses, but it does not automatically know which code files, database schema elements, requirements, policies, and access-control rules are connected. UCE turns those artifacts into an explicit graph and exposes deterministic reasoning through MCP tools. The benchmark evaluates whether that structure improves change-impact analysis and governance safety.

Across 24 scenarios, MCP-UCE raises file-impact F1 from 0.632 to 0.791. That improvement is not just cosmetic: file recall increases from 0.519 to 0.654, while file precision reaches 1.000. In plain terms, UCE finds more of the files that should be considered, and in this benchmark it does so without adding false positive impacted files.

Governance results are more nuanced. Requirement F1 improves from 0.416 to 0.486, and policy F1 improves from 0.521 to 0.603. The biggest governance improvement is precision: vanilla generates 25 aggregate governance false alarms, while MCP-UCE generates 0. However, both systems still miss many true governance violations, so the honest conclusion is that UCE improves trustworthiness but does not fully solve governance recall.

The RBAC result is the clearest safety signal. Under the repository RBAC policy, there are 34 denied write/delete probes. Vanilla permits 34 of them because it has no enforcement gate. MCP-UCE permits 0. This does not prove production security, but it does show the value of placing an explicit role-aware gate before mutation-capable tools.

The main tradeoff is latency and risk calibration. MCP-UCE is slower because it performs graph queries and constructs traces. It also has better risk ordering but worse absolute risk MAE in this run, which means it ranks scenarios more correctly even though its score scale still needs calibration.

What This Report Claims and Does Not Claim

- Claim: MCP-UCE improves file-impact F1, requirement F1, policy F1, and risk rank correlation on the generated 24-scenario benchmark.
- Claim: MCP-UCE eliminates simulated RBAC breaches on oracle-denied probes because it evaluates explicit RBAC rules before action execution.
- Claim: UCE reduces governance false alarms in this benchmark because graph-derived governance links are more precise than lexical guesses.
- Not claimed: this is not a full ICML-scale multi-repository evaluation.
- Not claimed: the vanilla baseline is not a fully optimized frontier-model tool-calling baseline.
- Not claimed: RBAC simulation is not a live penetration test; it is an enforcement comparison against encoded policy rules.

System Architecture

UCE is best understood as a control layer between an AI agent and the software project it wants to modify. The agent can still reason in natural language, but UCE prevents important context from living only inside model attention. Instead, UCE converts repository artifacts into graph entities and relationships, then serves deterministic MCP tools that answer questions like: which files depend on this schema element, which requirements are implicated, which policies enforce those requirements, and whether this user role is allowed to mutate this path.

Runtime Architecture Walkthrough




Figure 1. User-provided UCE runtime architecture.

The runtime diagram shows three important boundaries. First, identity comes from Keycloak, which attaches a role-bearing token to MCP calls. Second, the MCP gateway normalizes requests and becomes the controlled entry point for reasoning or action tools. Third, mutation-capable action tools are placed behind an RBAC guard. This placement matters because a direct LLM-to-database or direct LLM-to-filesystem path would skip the policy checks entirely.

A typical request lifecycle is: the agent receives a change request, calls a UCE reasoning tool, UCE checks the token, queries the graph, returns a trace, and only then can action tools be invoked. If the action would write or delete a governed path, the RBAC guard evaluates the rule set. Denies win over allows, more specific path patterns are prioritized, and role rank determines whether the principal satisfies the rule. The goal is not to make the LLM smarter by itself; the goal is to make the environment around the LLM less permissive and more auditable.

This is why the architecture is materially different from a vanilla LLM. A vanilla agent may retrieve snippets or infer likely impacted files, but it does not inherently know that editing `src/requirements` or `src/rbac` should require admin authority. UCE encodes that relationship as executable policy.

Ingestion Architecture Walkthrough




Figure 2. User-provided deterministic vs LLM ingestion architecture.

The ingestion diagram separates deterministic ingestion from optional LLM extraction. That separation is important. Deterministic parsers are appropriate for artifacts with strong structure: code, imports, function declarations, database schema, path patterns, and markdown frontmatter fields. LLM extraction is better treated as optional because it can improve recall on messy natural-language documents, but it also introduces policy-poisoning risk if untrusted text is merged into an authority graph without validation.

For this benchmark, the governance facts used in evaluation are structured: 10 requirement documents, 4 policy documents, and 7 RBAC rules. That makes deterministic parsing the right baseline. The LLM ingestion lane remains architecturally useful, but its outputs should be normalized, confidence-filtered, and reviewed before they become enforceable graph facts.

Governance Corpus

The project governance corpus is small but meaningful. Requirements describe schema-level obligations, policies group and enforce requirements, and RBAC rules constrain who can write or delete important paths. This section

explains what those artifacts contain, because the benchmark scores only make sense if the reader understands what is being traced.

Requirements

ID	Title	Description
RQ-001	Store core meeting status and timestamps	Each record in meetings must include status, created_at, and updated_at. These fields are required to track meeting lifecycle and audit updates.
RQ-002	Persist meeting recordings and transcripts	meetings must store recording_url and transcript_url so recorded artifacts remain accessible. This ensures recordings are linked to each meeting.
RQ-003	Meetings must link to user and agent	Each meetings row must include user_id and agent_id to bind the meeting to ownership and the assigned agent. Missing links are not allowed.
RQ-004	Agents require a name and instructions	Every agents record must store name and instructions. These fields define the agent identity and behavior.
RQ-005	Users must have verified contact data	Each user must have email and email_verified tracked for authentication. The email field is required for account recovery.
RQ-006	Sessions require tokens and expiry	Each session must store token and expires_at for access control. These fields are mandatory for session management.
RQ-007	Account credentials must be stored securely	The account table must include access_token and refresh_token for OAuth flows. These fields are required for re-authentication.
RQ-008	Verification records must be complete	Each verification entry must include identifier, value, and expires_at. These fields support secure verification workflows.
RQ-009	Store meeting summaries and timing	meetings must include summary along with started_at and ended_at to support duration reporting. These fields power post-meeting analytics.
RQ-010	Meetings must be titled for search	The meetings table must include name to support search and discovery. Meeting name is required for listing views.

These requirements are not abstract labels. They encode concrete obligations such as retaining meeting timestamps, preserving OAuth tokens, requiring session expiry fields, and ensuring verification records remain

complete. If an agent changes a table or column related to these obligations, the correct behavior is to surface the associated requirement before mutation.

Policies

Policy	Description	Enforces
P-001	Meeting data must preserve audit integrity.	RQ-001, RQ-003
P-002	Recorded meeting artifacts must remain available.	RQ-002, RQ-009
P-003	User identity and sessions must remain verifiable.	RQ-005, RQ-006
P-004	Agents must retain identity and configuration fields.	RQ-004, RQ-010

Policies are higher-level governance groupings. For example, P-001 links meeting audit integrity to requirements about meeting status, timestamps, ownership, and assigned agent. In evaluation, policy detection depends on whether the system can recover the requirements first and then traverse policy-enforcement links.

RBAC Rules

- DENY_NON_ADMIN_RBAC_WRITE
- DENY_NON_ADMIN_RBAC_DELETE
- DENY_NON_ADMIN_POLICIES_WRITE
- DENY_NON_ADMIN_REQUIREMENTS_WRITE
- ALLOW_EDITOR_WRITE_SRC
- ALLOW_ADMIN_WRITE_ANY
- ALLOW_ADMIN_DELETE_ANY

The RBAC policy has a sharp operational meaning. Non-admin users are denied writes to RBAC docs, policy docs, and requirement docs. Editors can write general src paths, but the more specific deny rules override broad allows for governance paths. Admins can write and delete broadly. This is exactly the kind of rule interaction a language model should not be asked to improvise. UCE evaluates it deterministically.

Benchmark Design

The benchmark has 24 scenarios: table-level schema changes, column-level changes, and file-level refactors. The graph snapshot records 6 tables, 52 columns, and 16 backend files indexed for comparison.

For each scenario, the evaluation compares predicted sets against oracle sets. The predicted sets are impacted files, violated requirements, and enforced policies. The oracle is graph-derived. This makes the benchmark especially appropriate for testing graph-faithful methods, but it also means the evaluation should be interpreted as measuring repository-structure recovery rather than broad programming ability.

Metric Definitions in Plain English

Precision answers: when the system says something is impacted, how often is it right?

Recall answers: out of everything truly impacted, how much did the system find?

F1 balances precision and recall. It is useful when a system can game one side by predicting either almost nothing or almost everything.

Risk MAE measures absolute numeric calibration error. Lower is better, but it does not tell whether the system ranks risky scenarios in the right order.

Risk Spearman correlation measures ranking quality. Higher means the system orders low-risk and high-risk scenarios more like the oracle.

RBAC breach rate measures how often a system would allow an action that policy says should be denied.

Main Results Explained

System	N	File Precision	File Recall	File F1	Requirement F1	Policy F1	Risk MAE	Risk Spearman	Latency ms
Vanilla LLM	24	0.806	0.519	0.632	0.416	0.521	24.83	0.356	2.92
MCP-UCE	24	1.000	0.654	0.791	0.486	0.603	27.67	0.647	14.90




Figure 3. Aggregate F1 scores for file, requirement, and policy prediction.

The file-impact result is the cleanest general improvement: F1 increases by 0.159. This happens because MCP-UCE preserves perfect precision while improving recall. Vanilla already avoids many false positives, but it misses many true impacted files. UCE still misses some, but fewer. In engineering terms, this means UCE gives the agent a better blast-radius estimate before code changes.

Requirement and policy scores improve, but they remain lower than file-impact scores. This is not surprising. Governance detection depends on two links: first from code or schema to requirements, then from requirements to policies. Any missing link reduces recall. UCE improves precision strongly because graph links prevent some lexical false alarms, but low recall shows that the current linkage rules are conservative.

Risk results split into calibration and ranking. MCP-UCE has higher risk MAE (27.67 vs 24.83), so its absolute score is less calibrated in this run. But its Spearman correlation is much higher (0.647 vs 0.356), which means it orders

scenarios more usefully for triage. For a pre-change review workflow, ranking can be operationally valuable even before absolute scores are perfectly calibrated.




Figure 4. Absolute risk calibration error.




Figure 5. Mean per-scenario latency.

Latency increases from 2.92 ms to 14.90 ms. The added cost is expected because UCE performs graph traversal, governance lookup, and trace construction. At this repository scale, the absolute overhead is still small enough for a safety check before mutation.

Entity-Type Results Explained

Entity	System	File F1	Req F1	Policy F1	Risk MAE	Latency ms
column	Vanilla LLM	0.413	0.636	0.639	10.33	2.80
column	MCP-UCE	1.000	0.917	0.833	11.92	16.85
file	Vanilla LLM	0.403	0.000	0.000	28.67	2.80
file	MCP-UCE	1.000	0.000	0.000	37.50	5.54
table	Vanilla LLM	0.594	0.447	0.578	50.00	3.29
table	MCP-UCE	0.594	0.464	0.578	49.33	20.35

Table-level scenarios are where UCE and vanilla are closest on file F1. That is because table names such as meetings, user, session, and account are lexically visible in code, so a simple baseline can often find many related files. UCE still adds governance structure, but the file-impact advantage is not dramatic here.

Column-level scenarios are where graph reasoning matters most. Column names such as agent_id, expires_at, email, or created_at can appear in narrow places, and lexical matching can miss transitive impact. MCP-UCE reaches file F1 of 1.000 for column scenarios, while vanilla is 0.413. This is the clearest evidence that explicit schema-to-code edges matter.

File-level scenarios also favor UCE because reverse import closure is a graph operation. Vanilla can guess files that mention similar tokens, but it cannot reliably recover dependency direction. MCP-UCE reaches file F1 of 1.000 for file scenarios, compared with vanilla at 0.403.




Figure 6. Scenario-level file F1 delta. Positive bars favor MCP-UCE.

Policy and Requirement Violation Analysis Explained

System	Req Caught	Req Missed	Req False	Policy Caught	Policy Missed	Policy False
--------	------------	------------	-----------	---------------	---------------	--------------

			Alarm			Alarm
Vanilla LLM	26	58	15	19	25	10
MCP-UCE	27	57	0	19	25	0

Governance Violation Capture (Higher Green, Lower Red)




Figure 7. Requirement and policy violation capture.

The requirement result is mixed but important. MCP-UCE catches 27 requirement violations while vanilla catches 26. That is only one additional true positive, so the recall improvement is modest. The bigger improvement is that MCP-UCE has 0 requirement false alarms while vanilla has 15. This means UCE is more trustworthy when it flags a requirement, but the system still needs stronger recall.

The policy result has identical true-positive and missed counts (19 caught and 25 missed for UCE; 19 caught and 25 missed for vanilla). The difference is false alarms: vanilla has 10, UCE has 0. This suggests that policy traversal from graph-linked requirements avoids speculative policy matches.

The correct interpretation is not that UCE fully solves policy and requirement violation detection. It does not. The correct interpretation is that UCE makes governance prediction cleaner and less noisy, while exposing where the graph needs richer semantic links to improve recall.

RBAC Breach Analysis Explained

System	Total Probes	Denied	Allowed	Breaches	Breach Rate	Blocked Denied	False Deny
Vanilla LLM	40	34	6	34	100.0%	0	0
MCP-UCE	40	34	6	0	0.0%	34	0




Figure 8. RBAC breach comparison on oracle-denied requests.

RBAC is where architecture matters more than model skill. A vanilla LLM baseline without an enforcement layer has no mechanism that forces it to deny a write or delete operation. It might say in prose that a change is risky, but unless the tool gateway enforces the rule, the action can still proceed. In this simulation, that means vanilla permits every oracle-denied request.

MCP-UCE blocks 34 denied requests and permits 0. This result follows directly from the RBAC guard: the action request is not treated as mere text, but as an operation over a path by a role. The guard evaluates that tuple against the rule graph.




Figure 9. RBAC breach rate by path group.

The path-group breakdown matters because governance files and backend files have different policy semantics. Governance paths such as `src/rbac`, `src/policies`, and `src/requirements` are deliberately stricter. Backend `src` paths can be editable by editors for write operations, but deletes remain admin-only in this policy. A good system must respect those distinctions at execution time.

Scenario-Level Interpretation

At the scenario level, MCP-UCE wins file F1 on 13 scenarios, ties on 11, and loses on 0. For requirement F1, it wins on 9, ties on 15, and loses on 0. For policy F1, it wins on 6, ties on 18, and loses on 0. This pattern says the file-impact component is much more mature than the governance-recall component.

Risk absolute error improves on 6 scenarios, ties on 7, and worsens on 11. That explains why the aggregate risk story is complicated: UCE ranks better overall, but the score formula still over- or under-shoots individual scenarios.




Figure 10. Scenario-level risk absolute error.

Design Implications

The benchmark points to three concrete design lessons. First, graph-backed impact analysis is useful immediately, especially for column and file changes. Second, governance precision benefits from deterministic policy links, but recall needs richer semantic modeling. Third, RBAC should be enforced outside the model as a tool-level gate. Asking the model to remember access-control policy is not enough.

The next engineering step should not be to make the prose prompt longer. It should be to improve graph coverage: better column-reference parsing, richer requirement-to-schema linking, requirement synonyms, policy expansion rules, and calibration data for the risk score. Those changes would attack the observed weaknesses directly.

Threats to Validity

- Single-repository scope: results may not generalize to large monorepos or different technology stacks.
- Oracle coupling: oracle labels are graph-derived, so graph-faithful methods have a structural advantage.
- Baseline limitation: vanilla is a no-tool lexical proxy, not a fully optimized frontier-model tool-calling agent.
- Postprocess caveat: the current artifact summary marks `postprocess_only=true`, so the latest document was generated from existing benchmark outputs and current RBAC docs rather than a fresh Neo4j rerun.
- RBAC simulation limitation: breach rates are calculated from policy evaluation over probes, not from live production attack traffic.

Conclusion

UCE is not just a context cache. It is a governance-aware control layer that changes what an agent can know and what it is allowed to do. The benchmark shows strong evidence for improved file-impact analysis, cleaner governance precision, better risk ranking, and explicit RBAC enforcement. The remaining weakness is governance recall and risk calibration, both of which are engineering problems that can be improved with better graph links and calibration data.

The honest bottom line is this: MCP-UCE is already meaningfully stronger than a no-tool vanilla baseline for governed change analysis, but the paper should present it as a promising system with clear next steps, not as a solved problem. That framing is accurate, credible, and much harder for a reviewer to dismiss.

Appendix A: Full Scenario Metrics

Scenario	Type	Entity	FileF 1 Vanil la	FileF 1 UCE	Delt a FileF 1	ReqF 1 Vanil la	ReqF 1 UCE	Delt a ReqF 1	PolF1 Vanil la	PolF 1 UCE	Delt a PolF 1	RiskE rr Vanil la	RiskE rr UCE	RiskE rr Delta UCE- Vanil la
COL-account-access_token	colu mn	account.access_token	1.000	1.00 0	0.00 0	0.667	1.00 0	0.33 3	0.000	0.00 0	0.00 0	7.000	7.000	0.000
COL-account-access_token_expires_at	colu mn	account.access_token_exp ires_at	1.000	1.00 0	0.00 0	0.000	0.00 0	0.00 0	0.000	0.00 0	0.00 0	11.00 0	7.000	- 4.000
COL-agents-created_at	colu mn	agents.created_at	0.000	1.00 0	1.00 0	0.667	1.00 0	0.33 3	0.667	1.00 0	0.33 3	7.000	7.000	0.000
COL-agents-id	colu mn	agents.id	0.952	1.00 0	0.04 8	1.000	1.00 0	0.00 0	1.000	1.00 0	0.00 0	2.000	66.00 0	64.00 0
COL-meetings-agent_id	colu mn	meetings.agent_id	0.000	1.00 0	1.00 0	0.400	1.00 0	0.60 0	0.500	1.00 0	0.50 0	26.00 0	7.000	- 19.00 0
COL-meetings-created_at	colu mn	meetings.created_at	0.000	1.00 0	1.00 0	0.400	1.00 0	0.60 0	0.500	1.00 0	0.50 0	26.00 0	7.000	- 19.00 0
COL-session-created_at	colu mn	session.created_at	0.000	1.00 0	1.00 0	0.667	1.00 0	0.33 3	0.667	1.00 0	0.33 3	5.000	7.000	2.000
COL-session-expires_at	colu mn	session.expires_at	0.000	1.00 0	1.00 0	1.000	1.00 0	0.00 0	1.000	1.00 0	0.00 0	2.000	7.000	5.000
COL-user-created_at	colu mn	user.created_at	0.000	1.00 0	1.00 0	0.500	1.00 0	0.50 0	0.667	1.00 0	0.33 3	19.00 0	7.000	- 12.00 0
COL-user-email	colu mn	user.email	0.000	1.00 0	1.00 0	0.667	1.00 0	0.33 3	0.667	1.00 0	0.33 3	15.00 0	7.000	- 8.000
COL-verification-created_at	colu mn	verification.created_at	1.000	1.00 0	0.00 0	0.667	1.00 0	0.33 3	1.000	1.00 0	0.00 0	4.000	7.000	3.000
COL-verification-expires_at	colu mn	verification.expires_at	1.000	1.00 0	0.00 0	1.000	1.00 0	0.00 0	1.000	1.00 0	0.00 0	0.000	7.000	7.000
FIL- src_app_api_auth_[...all]_r oute.ts	file	src/app/api/auth/[...all]/ro ute.ts	0.000	1.00 0	1.00 0	0.000	0.00 0	0.00 0	0.000	0.00 0	0.00 0	0.000	1.000	1.000
FIL- src_app_api_inngest_rout	file	src/app/api/inngest/route. ts	0.000	1.00 0	1.00 0	0.000	0.00 0	0.00 0	0.000	0.00 0	0.00 0	29.00 0	30.00 0	1.000

e.ts														
FIL-src_inngest_client.ts	file	src/inngest/client.ts	0.000	1.00 0	1.00 0	0.000	0.00 0	0.00 0	0.000	0.00 0	0.00 0	50.00 0	57.00 0	7.000
FIL-src_trpc_init.ts	file	src/trpc/init.ts	0.750	1.00 0	0.25 0	0.000	0.00 0	0.00 0	0.000	0.00 0	0.00 0	35.00 0	60.00 0	25.00 0
FIL-src_trpc_query-client.ts	file	src/trpc/query-client.ts	1.000	1.00 0	0.00 0	0.000	0.00 0	0.00 0	0.000	0.00 0	0.00 0	22.00 0	32.00 0	10.00 0
FIL-src_trpc_routers__app.ts	file	src/trpc/routers/_app.ts	0.667	1.00 0	0.33 3	0.000	0.00 0	0.00 0	0.000	0.00 0	0.00 0	36.00 0	45.00 0	9.000
TBL-account	table	account	0.429	0.42 9	0.00 0	0.500	0.50 0	0.00 0	0.500	0.50 0	0.00 0	61.00 0	61.00 0	0.000
TBL-agents	table	agents	0.842	0.84 2	0.00 0	0.333	0.33 3	0.00 0	0.667	0.66 7	0.00 0	30.00 0	30.00 0	0.000
TBL-meetings	table	meetings	0.818	0.81 8	0.00 0	0.727	0.83 3	0.10 6	1.000	1.00 0	0.00 0	27.00 0	23.00 0	- 4.000
TBL-session	table	session	0.533	0.53 3	0.00 0	0.286	0.28 6	0.00 0	0.500	0.50 0	0.00 0	66.00 0	66.00 0	0.000
TBL-user	table	user	0.778	0.77 8	0.00 0	0.500	0.50 0	0.00 0	0.800	0.80 0	0.00 0	32.00 0	32.00 0	0.000
TBL-verification	table	verification	0.167	0.16 7	0.00 0	0.333	0.33 3	0.00 0	0.000	0.00 0	0.00 0	84.00 0	84.00 0	0.000

Appendix B: RBAC Denied Probes

Probe	Path Group	Role	Operation	Path	System	Predicted Allowed	Breach	Decision Mode
RBAC-001	governance	viewer	write	src/rbac/RBAC_DEMO_001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-001	governance	viewer	write	src/rbac/RBAC_DEMO_001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-002	governance	viewer	write	src/policies/P-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-002	governance	viewer	write	src/policies/P-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-003	governance	viewer	write	src/requirements/RQ-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-003	governance	viewer	write	src/requirements/RQ-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-004	backend	viewer	write	src/app/api/auth/[...all]/route.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-004	backend	viewer	write	src/app/api/auth/[...all]/route.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-005	backend	viewer	write	src/app/api/inngest/route.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-005	backend	viewer	write	src/app/api/inngest/route.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-006	backend	viewer	write	src/inngest/client.ts	Vanilla LLM	True	1	no_rbac_guard

RBAC-006	backend	viewer	write	src/inngest/client.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-007	backend	viewer	write	src/trpc/init.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-007	backend	viewer	write	src/trpc/init.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-008	backend	viewer	write	src/trpc/query-client.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-008	backend	viewer	write	src/trpc/query-client.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-009	backend	viewer	write	src/trpc/routers/_app.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-009	backend	viewer	write	src/trpc/routers/_app.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-010	other	viewer	write	README.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-010	other	viewer	write	README.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-011	governance	viewer	delete	src/rbac/RBAC_DEMO_001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-011	governance	viewer	delete	src/rbac/RBAC_DEMO_001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-012	governance	viewer	delete	src/policies/P-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-012	governance	viewer	delete	src/policies/P-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-013	governance	viewer	delete	src/requirements/RQ-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-013	governance	viewer	delete	src/requirements/RQ-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-014	backend	viewer	delete	src/app/api/auth/[...all]/route.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-014	backend	viewer	delete	src/app/api/auth/[...all]/route.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-015	backend	viewer	delete	src/app/api/inngest/route.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-015	backend	viewer	delete	src/app/api/inngest/route.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-016	backend	viewer	delete	src/inngest/client.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-016	backend	viewer	delete	src/inngest/client.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-017	backend	viewer	delete	src/trpc/init.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-017	backend	viewer	delete	src/trpc/init.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-018	backend	viewer	delete	src/trpc/query-client.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-018	backend	viewer	delete	src/trpc/query-client.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-019	backend	viewer	delete	src/trpc/routers/_app.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-019	backend	viewer	delete	src/trpc/routers/_app.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-020	other	viewer	delete	README.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-020	other	viewer	delete	README.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-021	governance	editor	write	src/rbac/RBAC_DEMO_001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-021	governance	editor	write	src/rbac/RBAC_DEMO_001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-022	governance	editor	write	src/policies/P-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-022	governance	editor	write	src/policies/P-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-023	governance	editor	write	src/requirements/RQ-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-023	governance	editor	write	src/requirements/RQ-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-030	other	editor	write	README.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-030	other	editor	write	README.md	MCP-UCE	False	0	rbac_guard_enforced

RBAC-031	governance	editor	delete	src/rbac/RBAC_DEMO_001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-031	governance	editor	delete	src/rbac/RBAC_DEMO_001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-032	governance	editor	delete	src/policies/P-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-032	governance	editor	delete	src/policies/P-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-033	governance	editor	delete	src/requirements/RQ-001.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-033	governance	editor	delete	src/requirements/RQ-001.md	MCP-UCE	False	0	rbac_guard_enforced
RBAC-034	backend	editor	delete	src/app/api/auth/[...all]/route.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-034	backend	editor	delete	src/app/api/auth/[...all]/route.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-035	backend	editor	delete	src/app/api/inngest/route.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-035	backend	editor	delete	src/app/api/inngest/route.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-036	backend	editor	delete	src/inngest/client.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-036	backend	editor	delete	src/inngest/client.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-037	backend	editor	delete	src/trpc/init.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-037	backend	editor	delete	src/trpc/init.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-038	backend	editor	delete	src/trpc/query-client.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-038	backend	editor	delete	src/trpc/query-client.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-039	backend	editor	delete	src/trpc/routers/_app.ts	Vanilla LLM	True	1	no_rbac_guard
RBAC-039	backend	editor	delete	src/trpc/routers/_app.ts	MCP-UCE	False	0	rbac_guard_enforced
RBAC-040	other	editor	delete	README.md	Vanilla LLM	True	1	no_rbac_guard
RBAC-040	other	editor	delete	README.md	MCP-UCE	False	0	rbac_guard_enforced

References

1. Lewis et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks.
2. Yao et al. (2023). ReAct: Synergizing Reasoning and Acting in Language Models.
3. Schick et al. (2023). Toolformer: Language Models Can Teach Themselves to Use Tools.
4. Jimenez et al. (2024). SWE-bench: Can Language Models Resolve Real-World GitHub Issues?
5. Francis et al. (2018). Cypher: An Evolving Query Language for Property Graphs.
6. Musco et al. (2016). A Learning Algorithm for Change Impact Prediction.
7. Ji et al. (2023). Survey of Hallucination in Natural Language Generation.
8. Model Context Protocol (2024). Model Context Protocol Specification.
9. Tree-sitter maintainers (2026). tree-sitter incremental parsing system.

Extended Explanation Addendum

This addendum expands the argument behind the evaluation. It is intentionally prose-heavy: the goal is to make the report readable to someone who has not lived inside the codebase and needs to understand why the architecture exists, how to interpret the benchmark, and where the current evidence is strong or weak.

A. Why Vanilla LLM Context Is Not Enough

A vanilla LLM can reason over text that is present in its prompt or retrieved into its context window. That is powerful, but it creates a structural problem for governed software engineering: the important facts are often not co-located. A schema column may be referenced in one file, used indirectly through an imported function in another file, governed by a markdown requirement, grouped under a policy document, and protected by a path-based RBAC rule. The model can sometimes infer those links, but inference is not the same as enforcement or auditable traversal.

The failure mode is not that the model is unintelligent. The failure mode is that the system boundary is too soft. If a model sees a request like 'change session expiry handling,' it might identify the session table and nearby files. But unless the environment explicitly exposes dependency closure and governance relationships, the model has to guess which transitive files matter and which requirements must be considered. Those guesses can look confident even when they miss critical paths.

UCE changes the shape of the problem. Instead of asking the model to remember every artifact relationship, it puts those relationships into a graph and gives the model a controlled interface to query them. This turns hidden context into executable context. The important distinction is that the graph does not merely provide more text; it provides typed structure: File imports File, File references Column, Requirement governs Column, Policy enforces Requirement, AuthorityRule targets PathPattern, and so on.

That distinction explains why the file-impact result improves more than the governance result. File and import relationships are relatively concrete, so deterministic graph traversal helps immediately. Governance relationships are partly semantic, so the current graph still needs richer linking to catch all true violations. The benchmark result is therefore consistent with the architecture: the more structural the relation, the stronger UCE's immediate advantage.

B. Component-by-Component Runtime Explanation

The Keycloak identity plane gives each request an identity and role. In a governance system, this is not cosmetic. Role is part of the action tuple. The question is not only 'what file is being changed?' but 'who is trying to change it, with what operation, and under which policy?' A model response alone cannot provide a trustworthy answer to that question because the model is not the authority source for identity.

The MCP gateway is the entry point that turns agent intent into structured tool calls. This matters because natural language has ambiguity, but tool calls force the request into fields: operation, entity type, entity name, path, and role. Once the request is structured, deterministic checks can be applied before any mutation occurs.

The reasoning tools are read-oriented. They answer questions like `impact_analysis` and `risk_assessment` by traversing the graph. Because these tools are read-only, they are appropriate to expose broadly to the agent. They help the model understand consequences without letting it change the project.

The action tools are mutation-capable. They can write or delete, so they sit behind the RBAC guard. This separation is a core design principle. A system can be generous with read-only reasoning and strict with mutation. That is exactly what a human review workflow would do: let contributors inspect context, but require stronger authority to change policy or security-sensitive files.

The governance knowledge graph is the memory of the system. It stores project facts in a form that can be queried, checked, and audited. This is different from retrieval because retrieval returns passages, while the graph returns relationships. A returned path like Requirement -> Policy or Column -> File -> Importing File is evidence, not just prose.

The data plane, represented by Neo4j in this implementation, provides persistence and query semantics. Cypher queries can express transitive closure, policy linkage, and matching patterns more reliably than ad hoc prompt instructions. That is why the graph substrate is not a decorative implementation detail; it is central to reproducibility.

C. Component-by-Component Ingestion Explanation

The deterministic ingestion lane is responsible for artifacts where exact parsing is available. Source files, imports, functions, classes, database tables, columns, and structured markdown fields are all examples. The benefit of deterministic ingestion is repeatability: the same repository state should produce the same graph state.

The optional LLM ingestion lane exists for weaker structure, such as prose-heavy policy documents or requirements written without strict fields. The correct role for this lane is extraction assistance, not final authority. Any LLM-derived governance fact should be normalized, checked, and possibly reviewed before it becomes enforceable policy. This is why the diagram explicitly separates the lane and identifies poisoning risk.

Graph transport connects ingestion outputs to the persistent graph. This is where idempotence matters. A full refresh should merge or update nodes without duplicating entities, and deleted or changed files should not leave stale relationships behind. In a production deployment, this would be a major reliability concern because stale graph facts can mislead agents just as badly as missing facts.

The evaluation used deterministic governance parsing because the available governance docs were structured. This is an accuracy choice. It avoids claiming that an LLM extracted governance facts when the actual run had API quota limitations and relied on deterministic structured fields.

D. What the Benchmark Really Measures

The benchmark measures repository-grounded impact analysis. It does not measure general coding ability, patch correctness, or end-user product quality. Each scenario asks: given a table, column, or file change, can the system identify impacted files and linked governance obligations? This is a narrower question, but it is a valuable one because many agent failures happen before code generation: the agent edits the wrong scope or misses a policy implication.

The oracle is graph-derived. That is both a strength and a limitation. It is a strength because the labels are reproducible and tied to actual dependency structure. It is a limitation because a graph-based system is naturally aligned with a graph-derived oracle. A stronger future benchmark would combine graph-derived labels with human review labels, especially for governance semantics.

The vanilla baseline is intentionally conservative: it approximates what a no-tool model might do by using lexical relevance. This baseline is not the strongest possible LLM setup. A frontier model with carefully engineered prompts, retrieval, multiple passes, and tool use could do better. Therefore, the report should frame results as evidence that MCP-UCE beats a no-tool baseline, not as proof that it beats every possible LLM agent configuration.

The RBAC benchmark is a policy-enforcement simulation. It is not a live exploit test. The value of the result is architectural: without an enforcement gate, denied actions are permitted; with the gate, they are blocked. That is a meaningful comparison because access control is supposed to be enforced by the system, not merely recommended by the model.

E. Why the Results Look the Way They Do

The file-impact results improve because file impact is highly structural. Imports, references, table usage, and column usage form a graph. Vanilla lexical matching can find direct textual overlap, but it struggles with transitive dependency closure. UCE directly traverses that closure, so it recovers impacted files that are not obvious from token matching alone.

The table scenarios show smaller differences because table names are broad and often appear throughout the repository. Lexical matching has more signal in those cases. This is why both systems have similar table-level file F1. It does not mean UCE adds no value; it means the baseline is less disadvantaged when the entity name is globally visible.

The column scenarios are the strongest case for UCE. Columns are narrower, often appear in fewer places, and can be accessed through layers of code. If the graph knows that a file references a column, and another file imports that file, UCE can identify the transitive impact. Vanilla token matching can miss this because the importing file may not mention the column directly.

The file scenarios also favor UCE because reverse-import analysis is a graph query. If file A changes, the relevant downstream files are not necessarily files that share words with A. They are files that import A or depend on things A declares. That relationship is structural, not lexical.

The governance false-alarm improvement happens because UCE uses explicit links from requirements to policies. Vanilla can over-match policy text because it sees related words. UCE is stricter: if a requirement is not linked, the policy is not inferred. That reduces false positives, but it can also reduce recall when links are incomplete.

The risk result is split because risk is a derived score, not a directly observed label. UCE may identify a better relative ordering of scenarios while still using a formula that does not perfectly match the oracle scale. This is why Spearman improves while MAE worsens. The fix is calibration: tune the risk formula using validation data, or learn weights from labeled incidents while keeping the evidence trace deterministic.

F. Reviewer-Style Questions and Answers

- Q: Is this just RAG with a graph? A: No. RAG retrieves evidence text, while UCE encodes typed relationships and enforces RBAC before mutation. Retrieval can inform the model; UCE can also constrain action.
- Q: Does UCE replace the LLM? A: No. UCE narrows and grounds the LLM's reasoning. The model still interprets tasks and writes code, but UCE supplies deterministic context and policy gates.
- Q: Why is requirement recall still low? A: The current governance linker is conservative and relies on structured/documented links. It avoids false positives but misses semantic associations that richer extraction could recover.
- Q: Is the RBAC result too easy? A: It is easy in the sense that a system without an RBAC gate cannot enforce RBAC. That is exactly the architectural point. Access control must live in the tool boundary, not inside model intention.
- Q: Why keep the vanilla baseline if it is not a frontier model? A: It establishes a no-tool control condition. Future work should add stronger baselines, but the no-tool baseline is still useful for measuring the value of deterministic tooling.
- Q: What would make this paper stronger? A: A fresh full benchmark run, more repositories, human-reviewed oracles, live LLM tool-calling baselines, and a calibrated risk model.

G. Recommended Next Experiments

The first next experiment should be a fresh end-to-end run with Neo4j reachable, recording graph ingestion logs, scenario generation, prediction outputs, and exact environment metadata. This would remove the current postprocess-only caveat from the report.

The second experiment should add a live frontier-model baseline with the same token budget and no MCP tools, then a frontier-model baseline with generic retrieval, then MCP-UCE. That would separate the effect of better prompting from the effect of deterministic graph tooling.

The third experiment should expand governance recall. Add synonyms and semantic labels to requirements, test LLM-assisted extraction in a quarantined lane, and compare deterministic-only vs LLM-assisted graph construction. This would directly address the missed requirement and policy violations.

The fourth experiment should calibrate risk. Use scenario labels or human severity judgments to learn weights for affected files, affected functions, violated requirements, and enforced policies. Then report both rank correlation and absolute calibration.

H. Final Framing for a Strong Paper

The strongest honest framing is: UCE is a deterministic governance layer for software agents that improves structural impact analysis and enforces RBAC at the tool boundary. The current evidence is strongest for file impact and RBAC enforcement, promising but incomplete for governance recall, and mixed for risk calibration. That framing is credible because it acknowledges limitations while showing clear technical value.

The paper should not pretend that UCE has solved all agentic software safety. It should argue that agent safety improves when context and authority are moved out of the model's hidden reasoning and into explicit, queryable, enforceable infrastructure. That is the real contribution.